

**Department of Electronic and  
Telecommunication Engineering University  
of Moratuwa**

**EN3150 - Pattern Recognition**

**Assignment 03**

Simple convolutional neural network to perform  
classification



210141U	DISSANAYAKA D.M.S.P.
210341H	LIYANAARACHCHI L.A.S.
210303U	KULASINGHAM P.N.
210705E	WICKRAMASINGHA M.P.D.N.

Submission Date: 12.12.2024

# Contents

<b>1</b>	<b>CNN for Image Classification</b>	<b>3</b>
1.1	Prepare the Environment . . . . .	3
1.2	Add Dataset . . . . .	3
1.3	Split the dataset into training, validation, and testing . . . . .	3
1.4	Build the CNN Model . . . . .	4
1.4.1	Check Cuda Availability and Version . . . . .	5
1.5	Custom CNN Model Architecture . . . . .	6
1.6	Activation Function Selection . . . . .	7
1.6.1	ReLU (Rectified Linear Unit) . . . . .	7
1.6.2	Softmax (used via log_softmax) . . . . .	7
1.7	Train the Model . . . . .	7
1.8	Adam Optimizer over SGD . . . . .	10
1.9	Sparse Categorical Crossentropy as the Loss Function . . . . .	11
1.10	Testing the model . . . . .	11
1.11	Plot training and validation loss for 0.0001, 0.001, 0.01, and 0.1 . . . . .	13
<b>2</b>	<b>Comapre the Network with state-of-the-art Networks</b>	<b>23</b>
2.1	Dataset Overview . . . . .	23
2.2	ResNet . . . . .	24
2.2.1	Loading the pretrained model . . . . .	24
2.2.2	Fine-tuning the model . . . . .	24
2.2.3	Evaluating the model . . . . .	26
2.3	DenseNet . . . . .	30
2.3.1	Loading the pretrained model . . . . .	30
2.3.2	Fine tuning the model . . . . .	30
2.3.3	Evaluating the model . . . . .	32
2.4	Selecting pre-trained model or architecture . . . . .	35
2.5	Fine-tuning the pre-trained model . . . . .	35
2.6	Train the fine-tuned model . . . . .	36
2.7	Training and validation loss values . . . . .	36
2.8	Evaluate the fine-tuned model . . . . .	36
2.9	Compare the test accuracy . . . . .	36
2.10	Discussion trade-offs, advantages, and limitations . . . . .	36
2.10.1	Custom Model . . . . .	36
2.10.2	Pretrained Model . . . . .	37

# 1 CNN for Image Classification

## 1.1 Prepare the Environment

```
1 import torch
2 from torch.utils.data import DataLoader, Dataset
3 from torchvision import datasets, transforms
```

Listing 1: Python Code for Data Loading and Transformation

## 1.2 Add Dataset

```
1 dataset_path = './realwaste/realwaste-main/RealWaste'
```

Listing 2: Dataset Path Example

## 1.3 Split the dataset into training, validation, and testing

```
1 # Define transformations
2 transform = transforms.Compose([
3     transforms.Resize((128, 128)), # Resize images to 128x128
4     transforms.ToTensor(), # Convert images to PyTorch
5     tensors
6     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) #
7     Normalize (mean, std for RGB)
8 ])
9
10 # Load the dataset
11 dataset = datasets.ImageFolder(root=dataset_path, transform=
12     transform)
13
14 # Split into training and validation datasets
15 train_size = int(0.6 * len(dataset))
16 test_size = int(0.2 * len(dataset))
17 val_size = len(dataset) - train_size - test_size
18 train_dataset, test_dataset, val_dataset = torch.utils.data.
19     random_split(dataset, [train_size, test_size, val_size])
20
21 # Create DataLoaders
22 train_loader = DataLoader(train_dataset, batch_size=32,
23     shuffle=True)
24 test_loader = DataLoader(test_dataset, batch_size=32, shuffle=
25     False)
26 val_loader = DataLoader(val_dataset, batch_size=32, shuffle=
27     False)
28
29 # Print class names
30 print("Classes:", dataset.classes)
```

Listing 3: Dataset Loading and Transformation with PyTorch

Classes: ['Cardboard', 'Food Organics', 'Glass', 'Metal', 'Miscellaneous  
Trash', 'Paper', 'Plastic', 'Textile Trash', 'Vegetation']

## 1.4 Build the CNN Model

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5 class CustomCNN(nn.Module):
6     def __init__(self, x1, m1, x2, m2, x3, d, K):
7         super(CustomCNN, self).__init__()
8
9         # First Convolutional Layer
10        self.conv1 = nn.Conv2d(in_channels=3, out_channels=x1,
11                                kernel_size=m1, stride=1, padding=m1 // 2)
12        self.activation1 = nn.ReLU()
13
14        # MaxPooling Layer 1
15        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)
16
17        # Second Convolutional Layer
18        self.conv2 = nn.Conv2d(in_channels=x1, out_channels=x2,
19                                kernel_size=m2, stride=1, padding=m2 // 2)
20        self.activation2 = nn.ReLU()
21
22        # MaxPooling Layer 2
23        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)
24
25        # Fully Connected Layer
26        self.fc1 = nn.Linear(x2 * (128 // 4) * (128 // 4), x3)
27        # Adjust dimensions based on input size and
28        # pooling
29        self.activation3 = nn.ReLU()
30        self.dropout = nn.Dropout(d) # Dropout layer
31
32        # Output Layer
33        self.fc2 = nn.Linear(x3, K)
34
35    def forward(self, x):
36        x = self.conv1(x)
37        x = self.activation1(x)
38        x = self.pool1(x)
39
40        x = self.conv2(x)
41        x = self.activation2(x)
42        x = self.pool2(x)
43
44        x = torch.flatten(x, 1) # Flatten the output for the
45                                # fully connected layer
46
47        x = self.fc1(x)
48        x = self.activation3(x)
49        x = self.dropout(x)
50
51        x = self.fc2(x)
52        return F.log_softmax(x, dim=1) # Softmax activation
53
54    # Define the model parameters
55    x1, m1 = 32, 3 # Filters and kernel size for the first
56                   # convolutional layer
```

```

51 x2, m2 = 64, 3 # Filters and kernel size for the second
    convolutional layer
52 x3 = 128      # Number of units in the fully connected layer
53 d = 0.5       # Dropout rate
54 K = len(dataset.classes) # Number of output classes
55
56 # Instantiate the model
57 model = CustomCNN(x1, m1, x2, m2, x3, d, K)
58
59 # Print the model architecture
60 print(model)

```

Listing 4: Custom CNN Model in PyTorch

### CustomCNN Model Architecture:

**(conv1):** Conv2d(3, 32, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))  
**(activation1):** ReLU()  
**(pool1):** MaxPool2d(kernel\_size=2, stride=2, padding=0, dilation=1, ceil\_mode=False)  
**(conv2):** Conv2d(32, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))  
**(activation2):** ReLU()  
**(pool2):** MaxPool2d(kernel\_size=2, stride=2, padding=0, dilation=1, ceil\_mode=False)  
**(fc1):** Linear(in\_features=65536, out\_features=128, bias=True)  
**(activation3):** ReLU()  
**(dropout):** Dropout(p=0.5, inplace=False)  
**(fc2):** Linear(in\_features=128, out\_features=9, bias=True)

#### 1.4.1 Check Cuda Availability and Version

```

1 # Check if CUDA is available
2 cuda_available = torch.cuda.is_available()
3 print(f"CUDA available: {cuda_available}")
4
5 # If CUDA is available, check the PyTorch CUDA version
6 if cuda_available:
7     print(f"CUDA version supported by PyTorch: {torch.version.
8         cuda}")
9     print(f"Number of GPUs available: {torch.cuda.device_count
10         ()}")
11     print(f"Current GPU: {torch.cuda.get_device_name(0)}")
12 else:
13     print("CUDA is not available in this PyTorch installation.
14         ")

```

Listing 5: Check CUDA Availability in PyTorch

**CUDA available:** True  
**CUDA version supported by PyTorch:** 11.8  
**Number of GPUs available:** 1  
**Current GPU:** NVIDIA GeForce GTX 1650

## 1.5 Custom CNN Model Architecture

### **Conv1 Layer:**

*Input Channels:* 3 (RGB)

*Output Channels:* 32

*Kernel Size:* 3x3

*Stride:* 1

*Padding:* 1 (to maintain spatial dimensions)

*Activation Function:* ReLU

*Output Shape after Pooling:* (32, 64, 64)

### **Conv2 Layer:**

*Input Channels:* 32 (from Conv1)

*Output Channels:* 64

*Kernel Size:* 3x3

*Stride:* 1

*Padding:* 1 (to maintain spatial dimensions)

*Activation Function:* ReLU

*Output Shape after Pooling:* (64, 32, 32)

### **Flatten Layer:**

*Input Shape:* (64, 32, 32)

*Output Shape:* 65536 features

### **Fully Connected Layer (fc1):**

*Input Features:* 65536

*Output Units:* 128

*Activation Function:* ReLU

### **Dropout Layer:**

*Dropout Rate:* 0.5

### **Fully Connected Layer (fc2):**

*Input Features:* 128

*Output Classes:* K (number of output classes)

### **Softmax Output Layer:**

*Output:* Log probabilities over K classes

### **Number of Parameters Calculation:**

*Conv1 Parameters:*

Formula:  $(3 \times 3 \times 3 + 1) \times 32$

Calculation:  $(27 + 1) \times 32 = 896$

*Conv2 Parameters:*

Formula:  $(32 \times 3 \times 3 + 1) \times 64$

Calculation:  $(288 + 1) \times 64 = 18496$

*Fully Connected Layer 1 Parameters:*

Formula:  $(64 \times 32 \times 32 + 1) \times 128$   
Calculation:  $(65536 + 1) \times 128 = 8388608$

*Fully Connected Layer 2 Parameters:*

Formula:  $(128 + 1) \times K$   
Calculation:  $(129) \times K$

## 1.6 Activation Function Selection

### 1.6.1 ReLU (Rectified Linear Unit)

ReLU is chosen as the activation function for intermediate layers due to its efficiency and simplicity. It helps mitigate the vanishing gradient problem by maintaining non-zero gradients for positive values, enabling faster and more stable training. Additionally, ReLU introduces sparsity in the activations by setting negative values to zero, which:

- Enhances model interpretability.
- Helps reduce overfitting.

From a computational perspective, ReLU is simple and more efficient compared to other nonlinear activation functions like sigmoid or tanh, making it an ideal choice for deep learning models.

### 1.6.2 Softmax (used via `log_softmax`)

Softmax is applied in the output layer as it converts raw outputs (logits) into a probability distribution over  $K$  classes. This is essential for multi-class classification tasks where:

- Each output class must have a probability between 0 and 1.
- The predicted probabilities across all classes must sum to 1.

Softmax ensures compatibility with the `CrossEntropyLoss` function, which:

- Directly optimizes the log probabilities generated by `log_softmax`.
- Ensures proper gradient updates during backpropagation.

By using Softmax, the model can effectively handle classification problems with  $K$  classes.

## 1.7 Train the Model

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4
5 # Model, loss function, and optimizer
6 model = CustomCNN(x1, m1, x2, m2, x3, d, K) # Define your
    model
```

```

7 device = torch.device("cuda" if torch.cuda.is_available() else
8   "cpu")
9 model.to(device)
10 criterion = nn.CrossEntropyLoss() # Loss function for
11   classification
12 optimizer = optim.Adam(model.parameters(), lr=0.001) # Adam
13   optimizer
14 # Training parameters
15 num_epochs = 20 # Number of epochs
16 train_losses = []
17 val_losses = []
18 # Training and validation loop
19 for epoch in range(num_epochs):
20     # Training phase
21     model.train()
22     running_loss = 0.0
23     correct = 0
24     total = 0
25     for images, labels in train_loader:
26         images, labels = images.to(device), labels.to(device)
27
28         optimizer.zero_grad() # Clear gradients
29         outputs = model(images) # Forward pass
30         loss = criterion(outputs, labels) # Compute loss
31         loss.backward() # Backward pass
32         optimizer.step() # Update weights
33         running_loss += loss.item() * images.size(0)
34
35     # Train Accuracy
36     _, predicted = torch.max(outputs.data, 1)
37     total += labels.size(0)
38     correct += (predicted == labels).sum().item()
39
40     epoch_train_loss = running_loss / len(train_loader.dataset)
41     train_losses.append(epoch_train_loss)
42
43     # Calculate Train Accuracy as percentage
44     train_accuracy = 100 * correct / total
45
46     # Validation phase
47     model.eval()
48     val_loss = 0.0
49     correct = 0
50     total = 0
51     with torch.no_grad():
52         for images, labels in val_loader:
53             images, labels = images.to(device), labels.to(
54               device)
55             outputs = model(images)
56             loss = criterion(outputs, labels)
57             val_loss += loss.item() * images.size(0)
58
59     # Accuracy

```



```

60         _, predicted = torch.max(outputs, 1)
61         total += labels.size(0)
62         correct += (predicted == labels).sum().item()
63
64     epoch_val_loss = val_loss / len(val_loader.dataset)
65     val_losses.append(epoch_val_loss)
66     val_accuracy = 100 * correct / total
67
68     print(f"Epoch {epoch+1}/{num_epochs}, "
69           f"Train Loss: {epoch_train_loss:.4f}, "
70           f"Train Accuracy: {train_accuracy:.2f}%",
71           f"Validation Loss: {epoch_val_loss:.4f}, "
72           f"Validation Accuracy: {val_accuracy:.2f}%")

```

### Training and Validation Results

Epoch 1/20, Train Loss: 2.0285, Train Accuracy: 25.29% Validation Loss: 1.6530, Validation Accuracy: 48.05%

Epoch 2/20, Train Loss: 1.6442, Train Accuracy: 41.46% Validation Loss: 1.3348, Validation Accuracy: 55.21%

Epoch 3/20, Train Loss: 1.4046, Train Accuracy: 50.23% Validation Loss: 1.1384, Validation Accuracy: 60.36%

Epoch 4/20, Train Loss: 1.2585, Train Accuracy: 55.95% Validation Loss: 1.1247, Validation Accuracy: 58.36%

Epoch 5/20, Train Loss: 1.1186, Train Accuracy: 60.22% Validation Loss: 1.0748, Validation Accuracy: 61.93%

Epoch 6/20, Train Loss: 0.9956, Train Accuracy: 64.57% Validation Loss: 0.9856, Validation Accuracy: 63.83%

Epoch 7/20, Train Loss: 0.9213, Train Accuracy: 66.64% Validation Loss: 1.0158, Validation Accuracy: 62.46%

Epoch 8/20, Train Loss: 0.7986, Train Accuracy: 70.36% Validation Loss: 0.9215, Validation Accuracy: 67.72%

Epoch 9/20, Train Loss: 0.7113, Train Accuracy: 73.41% Validation Loss: 0.9926, Validation Accuracy: 63.72%

Epoch 10/20, Train Loss: 0.6452, Train Accuracy: 76.53% Validation Loss: 0.9566, Validation Accuracy: 67.82%

Epoch 11/20, Train Loss: 0.5382, Train Accuracy: 79.62% Validation Loss: 0.9583, Validation Accuracy: 66.46%

Epoch 12/20, Train Loss: 0.4669, Train Accuracy: 82.67% Validation Loss: 1.0385, Validation Accuracy: 67.30%

Epoch 13/20, Train Loss: 0.4692, Train Accuracy: 83.27% Validation Loss: 1.0108, Validation Accuracy: 69.19%

Epoch 14/20, Train Loss: 0.4084, Train Accuracy: 84.92% Validation Loss: 1.0913, Validation Accuracy: 66.67%

Epoch 15/20, Train Loss: 0.3217, Train Accuracy: 87.65% Validation Loss: 1.0809, Validation Accuracy: 67.40%

Epoch 16/20, Train Loss: 0.3394, Train Accuracy: 88.71% Validation Loss: 1.1911, Validation Accuracy: 65.83%

Epoch 17/20, Train Loss: 0.2878, Train Accuracy: 89.55% Validation Loss: 1.2849, Validation Accuracy: 65.30%

Epoch 18/20, Train Loss: 0.2712, Train Accuracy: 90.04% Validation Loss: 1.1358, Validation Accuracy: 68.03%  
Epoch 19/20, Train Loss: 0.2313, Train Accuracy: 91.72% Validation Loss: 1.2599, Validation Accuracy: 67.09%  
Epoch 20/20, Train Loss: 0.2782, Train Accuracy: 89.93% Validation Loss: 1.3685, Validation Accuracy: 66.25%

```
1 import matplotlib.pyplot as plt
2
3 plt.plot(range(1, num_epochs+1), train_losses, label='Training
  Loss')
4 plt.plot(range(1, num_epochs+1), val_losses, label='Validation
  Loss')
5 plt.xlabel('Epochs')
6 plt.ylabel('Loss')
7 plt.legend()
8 plt.title(f'Training and Validation Loss')
9 plt.show()
```




Figure 1: Training and Validation Loss

## 1.8 Adam Optimizer over SGD

The Adam optimizer is chosen because it combines the benefits of Momentum and RMSProp, making it well-suited for deep learning tasks. It dynamically adjusts the learning rate for each parameter using estimates of the first and second moments of gradients. This enables:

- Faster convergence.
- Better handling of sparse gradients.

In contrast, SGD (Stochastic Gradient Descent) with a fixed learning rate may require more manual tuning and tends to converge slower without additional techniques like momentum.

For tasks like classification, where the dataset may have varying complexity, Adam provides better efficiency and performance out of the box.

## 1.9 Sparse Categorical Crossentropy as the Loss Function

CrossEntropyLoss is specifically designed for classification problems with mutually exclusive classes. It works by comparing the predicted probability distribution over classes with the true class labels, penalizing incorrect predictions. This makes it highly suitable for multi-class classification tasks like this one.

The loss function's mathematical formulation inherently aligns with softmax outputs, ensuring proper gradient updates for learning discriminative features.

### 1.10 Testing the model

```

1 from sklearn.metrics import confusion_matrix, precision_score,
  recall_score, accuracy_score
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4
5 # Evaluate the model on the testing dataset
6 model.eval()
7 test_loss = 0.0
8 correct = 0
9 total = 0
10 all_labels = []
11 all_preds = []
12
13 # Iterate over the test set
14 with torch.no_grad():
15     for images, labels in test_loader:
16         images, labels = images.to(device), labels.to(device)
17         outputs = model(images)
18         loss = criterion(outputs, labels)
19         test_loss += loss.item() * images.size(0)
20
21     # Get predictions
22     _, predicted = torch.max(outputs, 1)
23     total += labels.size(0)
24     correct += (predicted == labels).sum().item()
25
26     # Store all labels and predictions for confusion
27     # matrix
28     all_labels.extend(labels.cpu().numpy())
29     all_preds.extend(predicted.cpu().numpy())
30
31 # Calculate Test Accuracy

```

```

31 test_accuracy = 100 * correct / total
32 test_loss /= len(test_loader.dataset)
33
34 # Print test accuracy and loss
35 print(f"Test Loss: {test_loss:.4f}")
36 print(f"Test Accuracy: {test_accuracy:.2f}%")
37
38 # Confusion Matrix
39 cm = confusion_matrix(all_labels, all_preds)
40
41 # Plot confusion matrix
42 plt.figure(figsize=(8, 6))
43 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels
44             =range(K), yticklabels=range(K))
45 plt.title('Confusion Matrix')
46 plt.xlabel('Predicted Label')
47 plt.ylabel('True Label')
48 plt.show()
49
50 # Precision and Recall
51 precision = precision_score(all_labels, all_preds, average='
52 weighted')
53 recall = recall_score(all_labels, all_preds, average='weighted
54 ')
55
56 print(f"Precision (Weighted): {precision:.4f}")
57 print(f"Recall (Weighted): {recall:.4f}")
58
59 # Evaluate train accuracy
60 model.train()
61 train_loss = 0.0
62 correct = 0
63 total = 0
64 with torch.no_grad():
65     for images, labels in train_loader:
66         images, labels = images.to(device), labels.to(device)
67         outputs = model(images)
68         loss = criterion(outputs, labels)
69         train_loss += loss.item() * images.size(0)
70
71         _, predicted = torch.max(outputs.data, 1)
72         total += labels.size(0)
73         correct += (predicted == labels).sum().item()
74
75 train_accuracy = 100 * correct / total
76 print(f"Train Accuracy: {train_accuracy:.2f}%")

```

Listing 6: Model Evaluation and Metrics

Test Loss	1.4439
Test Accuracy	65.47%
Precision (Weighted)	0.6611
Recall (Weighted)	0.6547
Train Accuracy	91.79%

Table 1: Test and Train Performance Metrics




Figure 2: Confusion matrix

### 1.11 Plot training and validation loss for 0.0001, 0.001, 0.01, and 0.1

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import matplotlib.pyplot as plt
5
6 # Model, loss function, and optimizer
7 model = CustomCNN(x1, m1, x2, m2, x3, d, K) # Define your
8 model
9 device = torch.device("cuda" if torch.cuda.is_available() else
10 "cpu")
11 model.to(device)
12
13 LR = [0.0001, 0.001, 0.01, 0.1]
14
15 for lr in LR:
16     print("-----")
17     print(f"Learning Rate: {lr}")
18     criterion = nn.CrossEntropyLoss() # Loss function for
19     classification

```

```

17     optimizer = optim.Adam(model.parameters(), lr=lr) # Adam
18     optimizer
19
19     # Training parameters
20     num_epochs = 20 # Number of epochs
21     train_losses = []
22     val_losses = []
23
24     # Training and validation loop
25     for epoch in range(num_epochs):
26         # Training phase
27         model.train()
28         running_loss = 0.0
29         correct = 0
30         total = 0
31         for images, labels in train_loader:
32             images, labels = images.to(device), labels.to(
33                 device)
34
35             optimizer.zero_grad() # Clear gradients
36             outputs = model(images) # Forward pass
37             loss = criterion(outputs, labels) # Compute loss
38             loss.backward() # Backward pass
39             optimizer.step() # Update weights
40
41             running_loss += loss.item() * images.size(0)
42
43             # Train Accuracy
44             _, predicted = torch.max(outputs.data, 1)
45             total += labels.size(0)
46             correct += (predicted == labels).sum().item()
47
48         epoch_train_loss = running_loss / len(train_loader.
49             dataset)
50         train_losses.append(epoch_train_loss)
51
52         # Calculate Train Accuracy as percentage
53         train_accuracy = 100 * correct / total
54
55         # Validation phase
56         model.eval()
57         val_loss = 0.0
58         correct = 0
59         total = 0
60         with torch.no_grad():
61             for images, labels in val_loader:
62                 images, labels = images.to(device), labels.to(
63                     device)
64
65                 outputs = model(images)
66                 loss = criterion(outputs, labels)
67                 val_loss += loss.item() * images.size(0)
68
69                 # Accuracy
70                 _, predicted = torch.max(outputs, 1)
71                 total += labels.size(0)
72                 correct += (predicted == labels).sum().item()
73
74         epoch_val_loss = val_loss / len(val_loader.dataset)

```

```

71     val_losses.append(epoch_val_loss)
72     val_accuracy = 100 * correct / total
73
74     print(f"Epoch {epoch+1}/{num_epochs}, "
75           f"Train Loss: {epoch_train_loss:.4f}, "
76           f"Train Accuracy: {train_accuracy:.2f}% ",
77           f"Validation Loss: {epoch_val_loss:.4f}, "
78           f"Validation Accuracy: {val_accuracy:.2f}%")
79
80     print(f"Testing for Learning Rate: {lr}")
81     model.eval()
82     test_loss = 0.0
83     correct = 0
84     total = 0
85     with torch.no_grad():
86         for images, labels in test_loader:
87             images, labels = images.to(device), labels.to(
88                 device)
89             outputs = model(images)
90             loss = criterion(outputs, labels)
91             test_loss += loss.item() * images.size(0)
92
93             _, predicted = torch.max(outputs, 1)
94             total += labels.size(0)
95             correct += (predicted == labels).sum().item()
96
97     test_loss /= len(test_loader.dataset)
98     test_accuracy = 100 * correct / total
99
100    print(f"Learning Rate: {lr} | Test Loss: {test_loss:.4f},
101          Test Accuracy: {test_accuracy:.2f}%")
102
103    plt.plot(range(1, num_epochs+1), train_losses, label='
104              Training Loss')
105    plt.plot(range(1, num_epochs+1), val_losses, label='
106              Validation Loss')
107    plt.xlabel('Epochs')
108    plt.ylabel('Loss')
109    plt.legend()
110    plt.title(f'Training and Validation Loss for Learning Rate
111              : {lr}')
112    plt.show()

```

Listing 7: Training and Testing with Different Learning Rates

---

Learning Rate: 0.0001

Epoch 1/20, Train Loss: 2.0171, Train Accuracy: 27.46% Validation Loss: 1.7636, Validation Accuracy: 44.58%

Epoch 2/20, Train Loss: 1.7602, Train Accuracy: 37.60% Validation Loss: 1.5222, Validation Accuracy: 49.42%

Epoch 3/20, Train Loss: 1.6039, Train Accuracy: 44.16% Validation Loss: 1.4079, Validation Accuracy: 54.47%

Epoch 4/20, Train Loss: 1.5052, Train Accuracy: 47.14% Validation Loss: 1.3314, Validation Accuracy: 56.57%

Epoch 5/20, Train Loss: 1.4239, Train Accuracy: 50.68% Validation Loss:

1.2581, Validation Accuracy: 56.47%  
Epoch 6/20, Train Loss: 1.3516, Train Accuracy: 52.79% Validation Loss:  
1.2153, Validation Accuracy: 57.73%  
Epoch 7/20, Train Loss: 1.2865, Train Accuracy: 55.98% Validation Loss:  
1.1827, Validation Accuracy: 58.99%  
Epoch 8/20, Train Loss: 1.2529, Train Accuracy: 56.09% Validation Loss:  
1.1136, Validation Accuracy: 62.46%  
Epoch 9/20, Train Loss: 1.1713, Train Accuracy: 59.42% Validation Loss:  
1.0845, Validation Accuracy: 62.04%  
Epoch 10/20, Train Loss: 1.1410, Train Accuracy: 60.19% Validation Loss:  
1.0591, Validation Accuracy: 63.30%  
Epoch 11/20, Train Loss: 1.0907, Train Accuracy: 60.96% Validation Loss:  
1.0313, Validation Accuracy: 64.56%  
Epoch 12/20, Train Loss: 1.0682, Train Accuracy: 62.93% Validation Loss:  
1.0005, Validation Accuracy: 64.98%  
Epoch 13/20, Train Loss: 1.0176, Train Accuracy: 63.77% Validation Loss:  
1.0000, Validation Accuracy: 64.67%  
Epoch 14/20, Train Loss: 0.9746, Train Accuracy: 66.26% Validation Loss:  
0.9547, Validation Accuracy: 65.93%  
Epoch 15/20, Train Loss: 0.9785, Train Accuracy: 64.22% Validation Loss:  
0.9440, Validation Accuracy: 66.35%  
Epoch 16/20, Train Loss: 0.9388, Train Accuracy: 67.98% Validation Loss:  
0.9565, Validation Accuracy: 66.46%  
Epoch 17/20, Train Loss: 0.8845, Train Accuracy: 69.87% Validation Loss:  
0.9258, Validation Accuracy: 66.25%  
Epoch 18/20, Train Loss: 0.8450, Train Accuracy: 70.57% Validation Loss:  
0.9274, Validation Accuracy: 67.51%  
Epoch 19/20, Train Loss: 0.8313, Train Accuracy: 69.69% Validation Loss:  
0.9170, Validation Accuracy: 68.56%  
Epoch 20/20, Train Loss: 0.8069, Train Accuracy: 70.99% Validation Loss:  
0.9138, Validation Accuracy: 69.30%

**Testing for Learning Rate: 0.0001**

Learning Rate: 0.0001 — Test Loss: 0.9721, Test Accuracy: 65.16%



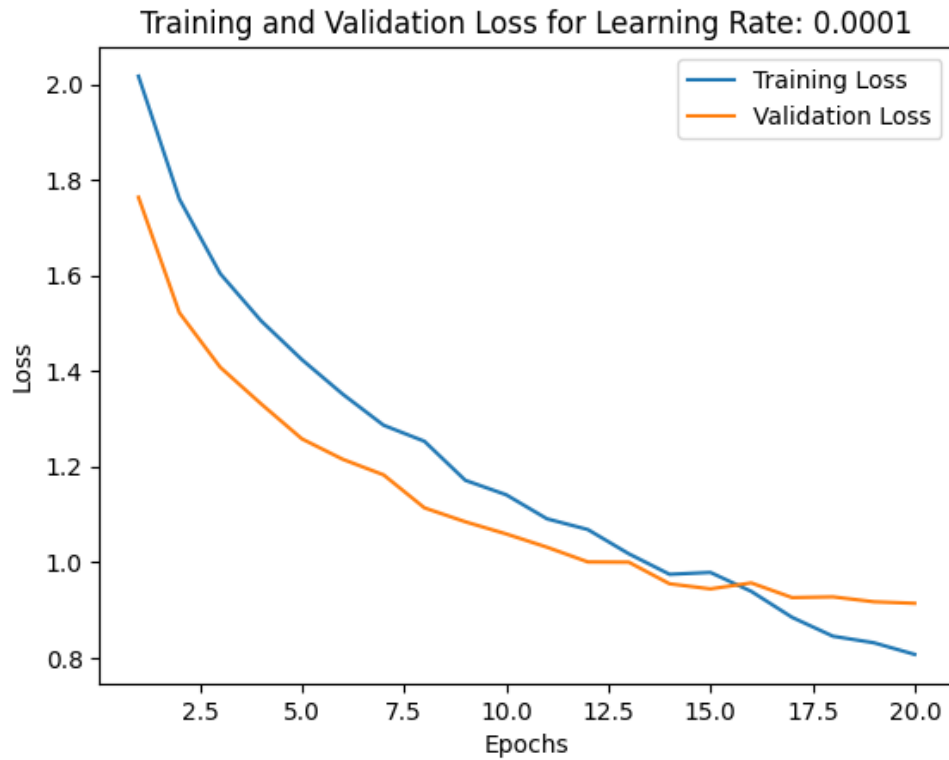


Figure 3: Learning Rate: 0.0001

---

Learning Rate: 0.001

Epoch 1/20, Train Loss: 1.3445, Train Accuracy: 51.56% Validation Loss: 1.0834, Validation Accuracy: 62.04%

Epoch 2/20, Train Loss: 1.1354, Train Accuracy: 58.82% Validation Loss: 1.0473, Validation Accuracy: 63.41%

Epoch 3/20, Train Loss: 1.0066, Train Accuracy: 64.43% Validation Loss: 1.0251, Validation Accuracy: 63.51%

Epoch 4/20, Train Loss: 0.9036, Train Accuracy: 67.66% Validation Loss: 1.0008, Validation Accuracy: 65.09%

Epoch 5/20, Train Loss: 0.8048, Train Accuracy: 70.26% Validation Loss: 0.9781, Validation Accuracy: 66.35%

Epoch 6/20, Train Loss: 0.7116, Train Accuracy: 74.39% Validation Loss: 0.9636, Validation Accuracy: 66.98%

Epoch 7/20, Train Loss: 0.6423, Train Accuracy: 76.01% Validation Loss: 1.0547, Validation Accuracy: 65.93%

Epoch 8/20, Train Loss: 0.5508, Train Accuracy: 79.97% Validation Loss: 1.0815, Validation Accuracy: 65.51%

Epoch 9/20, Train Loss: 0.5248, Train Accuracy: 80.88% Validation Loss: 1.0207, Validation Accuracy: 67.51%

Epoch 10/20, Train Loss: 0.4397, Train Accuracy: 83.58% Validation Loss: 1.1289, Validation Accuracy: 66.46%

Epoch 11/20, Train Loss: 0.4277, Train Accuracy: 83.30% Validation Loss:

1.1988, Validation Accuracy: 67.51%  
 Epoch 12/20, Train Loss: 0.4054, Train Accuracy: 84.95% Validation Loss: 1.1854, Validation Accuracy: 66.88%  
 Epoch 13/20, Train Loss: 0.3514, Train Accuracy: 86.92% Validation Loss: 1.1455, Validation Accuracy: 66.67%  
 Epoch 14/20, Train Loss: 0.3059, Train Accuracy: 88.53% Validation Loss: 1.2260, Validation Accuracy: 67.09%  
 Epoch 15/20, Train Loss: 0.2670, Train Accuracy: 89.55% Validation Loss: 1.4207, Validation Accuracy: 66.04%  
 Epoch 16/20, Train Loss: 0.2769, Train Accuracy: 89.41% Validation Loss: 1.2646, Validation Accuracy: 68.03%  
 Epoch 17/20, Train Loss: 0.2358, Train Accuracy: 91.93% Validation Loss: 1.3290, Validation Accuracy: 66.98%  
 Epoch 18/20, Train Loss: 0.2464, Train Accuracy: 90.60% Validation Loss: 1.3245, Validation Accuracy: 66.56%  
 Epoch 19/20, Train Loss: 0.2297, Train Accuracy: 90.88% Validation Loss: 1.3893, Validation Accuracy: 66.25%  
 Epoch 20/20, Train Loss: 0.2012, Train Accuracy: 92.00% Validation Loss: 1.4655, Validation Accuracy: 68.98%

#### Testing for Learning Rate: 0.001

Learning Rate: 0.001 — Test Loss: 1.4845, Test Accuracy: 66.63%




Figure 4: Learning Rate: 0.001

---

Learning Rate: 0.01

Epoch 1/20, Train Loss: 1.9594, Train Accuracy: 39.25% Validation Loss: 1.8079, Validation Accuracy: 33.96%  
Epoch 2/20, Train Loss: 1.4480, Train Accuracy: 51.84% Validation Loss: 1.5077, Validation Accuracy: 45.95%  
Epoch 3/20, Train Loss: 1.0856, Train Accuracy: 63.42% Validation Loss: 1.7044, Validation Accuracy: 48.90%  
Epoch 4/20, Train Loss: 0.8039, Train Accuracy: 73.73% Validation Loss: 1.7050, Validation Accuracy: 49.42%  
Epoch 5/20, Train Loss: 0.6031, Train Accuracy: 79.38% Validation Loss: 2.1100, Validation Accuracy: 41.11%  
Epoch 6/20, Train Loss: 0.5099, Train Accuracy: 83.30% Validation Loss: 2.4433, Validation Accuracy: 45.85%  
Epoch 7/20, Train Loss: 0.6892, Train Accuracy: 79.73% Validation Loss: 2.5181, Validation Accuracy: 41.22%  
Epoch 8/20, Train Loss: 0.5362, Train Accuracy: 83.48% Validation Loss: 2.4963, Validation Accuracy: 40.38%  
Epoch 9/20, Train Loss: 0.5140, Train Accuracy: 85.44% Validation Loss: 3.2499, Validation Accuracy: 37.54%  
Epoch 10/20, Train Loss: 0.6119, Train Accuracy: 82.32% Validation Loss: 3.2854, Validation Accuracy: 43.11%  
Epoch 11/20, Train Loss: 0.5154, Train Accuracy: 85.90% Validation Loss: 3.7873, Validation Accuracy: 42.27%  
Epoch 12/20, Train Loss: 0.5454, Train Accuracy: 86.46% Validation Loss: 2.8227, Validation Accuracy: 38.49%  
Epoch 13/20, Train Loss: 0.4634, Train Accuracy: 86.11% Validation Loss: 2.7984, Validation Accuracy: 43.11%  
Epoch 14/20, Train Loss: 0.3785, Train Accuracy: 88.74% Validation Loss: 3.0848, Validation Accuracy: 43.11%  
Epoch 15/20, Train Loss: 0.3061, Train Accuracy: 90.11% Validation Loss: 4.1324, Validation Accuracy: 40.38%  
Epoch 16/20, Train Loss: 0.4191, Train Accuracy: 89.37% Validation Loss: 3.9586, Validation Accuracy: 42.27%  
Epoch 17/20, Train Loss: 0.4112, Train Accuracy: 88.64% Validation Loss: 3.1025, Validation Accuracy: 40.80%  
Epoch 18/20, Train Loss: 0.3348, Train Accuracy: 90.56% Validation Loss: 3.5894, Validation Accuracy: 43.53%  
Epoch 19/20, Train Loss: 0.4569, Train Accuracy: 89.30% Validation Loss: 4.0916, Validation Accuracy: 40.17%  
Epoch 20/20, Train Loss: 0.4883, Train Accuracy: 86.43% Validation Loss: 3.4627, Validation Accuracy: 39.01%

**Testing for Learning Rate: 0.01**

Learning Rate: 0.01 — Test Loss: 3.2796, Test Accuracy: 41.79%




Figure 5: Learning Rate: 0.01

---

Learning Rate: 0.1

Epoch 1/20, Train Loss: 17.2835, Train Accuracy: 20.52% Validation Loss: 2.1676, Validation Accuracy: 17.14%

Epoch 2/20, Train Loss: 2.1520, Train Accuracy: 19.33% Validation Loss: 2.1776, Validation Accuracy: 17.14%

Epoch 3/20, Train Loss: 2.1473, Train Accuracy: 19.64% Validation Loss: 2.1865, Validation Accuracy: 17.14%

Epoch 4/20, Train Loss: 2.1508, Train Accuracy: 19.64% Validation Loss: 2.1825, Validation Accuracy: 17.14%

Epoch 5/20, Train Loss: 2.1469, Train Accuracy: 19.05% Validation Loss: 2.1679, Validation Accuracy: 17.14%

Epoch 6/20, Train Loss: 2.150

6, Train Accuracy: 18.77% Validation Loss: 2.1830, Validation Accuracy: 17.14% Epoch 7/20, Train Loss: 2.1479, Train Accuracy: 18.84% Validation Loss: 2.1785, Validation Accuracy: 17.14%

Epoch 8/20, Train Loss: 2.1466, Train Accuracy: 19.71% Validation Loss: 2.1761, Validation Accuracy: 16.09%

Epoch 9/20, Train Loss: 2.1510, Train Accuracy: 18.91% Validation Loss: 2.1902, Validation Accuracy: 17.14%

Epoch 10/20, Train Loss: 2.1473, Train Accuracy: 19.40% Validation Loss: 2.1805, Validation Accuracy: 17.14%

Epoch 11/20, Train Loss: 2.1508, Train Accuracy: 19.57% Validation Loss:

2.1717, Validation Accuracy: 17.14%  
Epoch 12/20, Train Loss: 2.1516, Train Accuracy: 19.08% Validation Loss:  
2.1907, Validation Accuracy: 17.14%  
Epoch 13/20, Train Loss: 2.1457, Train Accuracy: 20.24% Validation Loss:  
2.2028, Validation Accuracy: 17.14%  
Epoch 14/20, Train Loss: 2.1535, Train Accuracy: 18.73% Validation Loss:  
2.1788, Validation Accuracy: 17.14%  
Epoch 15/20, Train Loss: 2.1487, Train Accuracy: 19.96% Validation Loss:  
2.1807, Validation Accuracy: 17.14%  
Epoch 16/20, Train Loss: 2.1483, Train Accuracy: 18.84% Validation Loss:  
2.1727, Validation Accuracy: 17.14%  
Epoch 17/20, Train Loss: 2.1495, Train Accuracy: 19.36% Validation Loss:  
2.1702, Validation Accuracy: 17.14%  
Epoch 18/20, Train Loss: 2.1499, Train Accuracy: 19.47% Validation Loss:  
2.1704, Validation Accuracy: 17.14%  
Epoch 19/20, Train Loss: 2.1468, Train Accuracy: 20.10% Validation Loss:  
2.1739, Validation Accuracy: 17.14%  
Epoch 20/20, Train Loss: 2.1489, Train Accuracy: 19.78% Validation Loss:  
2.1723, Validation Accuracy: 17.14%

#### Testing for Learning Rate: 0.1

Learning Rate: 0.1 — Test Loss: 2.1360, Test Accuracy: 19.47%




Figure 6: Learning Rate: 0.1

### Comments on Results

- **Learning Rate = 0.0001**: Likely too slow; the loss decreases very gradually.
- **Learning Rate = 0.001**: Typically a good balance between speed and stability.
- **Learning Rate = 0.01**: Faster convergence but may start to oscillate.
- **Learning Rate = 0.1**: May diverge or show instability.

Based on the observations, **0.001** is selected as the optimal learning rate as it provides steady convergence of the validation loss without overfitting or oscillations.

## 2 Comapre the Network with state-of-the-art Networks

### 2.1 Dataset Overview

```
1 dataset_path = './realwaste/realwaste-main/RealWaste'
2
3 import torch
4 import torch.nn as nn
5 import torch.optim as optim
6 from torchvision import datasets, models, transforms
7 from torch.utils.data import DataLoader
8
9 num_classes = 9
10 batch_size = 32
11 learning_rate = 0.001
12 num_epochs = 20
13
14 from torchvision import transforms, datasets
15 from torch.utils.data import DataLoader, random_split
16
17 transform = transforms.Compose([
18     transforms.Resize((128, 128)),
19     transforms.ToTensor(),
20     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
21 ])
22
23 dataset = datasets.ImageFolder(root=dataset_path, transform=
    transform)
24 dataset_size = len(dataset)
25 indices = torch.randperm(dataset_size).tolist()
26
27 train_size = int(0.6 * dataset_size)
28 val_size = int(0.2 * dataset_size)
29 test_size = dataset_size - train_size - val_size
30
31 train_indices, val_indices, test_indices = indices[:train_size]
32     , indices[train_size:train_size+val_size], indices[
33     train_size+val_size:]
34
35 train_dataset = torch.utils.data.Subset(dataset, train_indices
36     )
37 val_dataset = torch.utils.data.Subset(dataset, val_indices)
38 test_dataset = torch.utils.data.Subset(dataset, test_indices)
39
40 # Print dataset information
41 print("Classes:", dataset.classes)
42 print(f"Total images: {len(dataset)}")
43 print(f"Training set size: {len(train_dataset)}")
44 print(f"Validation set size: {len(val_dataset)}")
45 print(f"Test set size: {len(test_dataset)}")
```

- **Classes:** Cardboard, Food Organics, Glass, Metal, Miscellaneous Trash, Paper, Plastic, Textile Trash, Vegetation

- Total images: 4752
- Training set size: 2851
- Validation set size: 950
- Test set size: 951

## 2.2 ResNet

### 2.2.1 Loading the pretrained model

```

1 model = models.resnet50(weights=models.ResNet50_Weights.
    IMAGENET1K_V1)
2
3 for param in model.parameters():
4     param.requires_grad = False
5
6 model.fc = nn.Linear(model.fc.in_features, num_classes)
7
8 nn.init.xavier_uniform_(model.fc.weight)
9 nn.init.zeros_(model.fc.bias)
10
11 device = torch.device("cuda" if torch.cuda.is_available() else
    "cpu")
12 model = model.to(device)
13
14 criterion = nn.CrossEntropyLoss()
15 optimizer = optim.Adam(model.fc.parameters(), lr=learning_rate
    )
16
17 print(f"Model is running on {device}")

```

Model is running on cuda

### 2.2.2 Fine-tuning the model

```

1 # Training and Validation
2 num_epochs = 20
3 train_losses = []
4 val_losses = []
5 train_accuracies = []
6 val_accuracies = []
7
8 for epoch in range(num_epochs):
9     # Training
10    model.train()
11    train_loss = 0.0
12    correct_train = 0
13    total_train = 0
14
15    for inputs, labels in train_loader:
16        inputs, labels = inputs.to(device), labels.to(device)
17
18        optimizer.zero_grad()
19        outputs = model(inputs)

```



```

20         loss = criterion(outputs, labels)
21         loss.backward()
22         optimizer.step()
23
24         train_loss += loss.item() * inputs.size(0)
25         _, predicted = outputs.max(1)
26         correct_train += (predicted == labels).sum().item()
27         total_train += labels.size(0)
28
29     train_accuracy = correct_train / total_train
30     train_loss /= total_train
31     train_losses.append(train_loss)
32     train_accuracies.append(train_accuracy)
33
34     # Validation
35     model.eval()
36     val_loss = 0.0
37     correct_val = 0
38     total_val = 0
39
40     with torch.no_grad():
41         for inputs, labels in val_loader:
42             inputs, labels = inputs.to(device), labels.to(
43                 device)
44
45             outputs = model(inputs)
46             loss = criterion(outputs, labels)
47
48             val_loss += loss.item() * inputs.size(0)
49             _, predicted = outputs.max(1)
50             correct_val += (predicted == labels).sum().item()
51             total_val += labels.size(0)
52
53     val_accuracy = correct_val / total_val
54     val_loss /= total_val
55     val_losses.append(val_loss)
56     val_accuracies.append(val_accuracy)
57
58     print(f"Epoch {epoch+1}/{num_epochs}, "
59           f"Train Loss: {train_loss:.4f}, "
60           f"Train Accuracy: {train_accuracy:.2f}%",
61           f"Validation Loss: {val_loss:.4f}, "
62           f"Validation Accuracy: {val_accuracy:.2f}%")

```

## Results

Epoch 1/20, Train Loss: 1.4362, Train Accuracy: 0.49% Validation Loss: 1.0811, Validation Accuracy: 0.61%

Epoch 2/20, Train Loss: 0.9367, Train Accuracy: 0.68% Validation Loss: 1.0090, Validation Accuracy: 0.64%

Epoch 3/20, Train Loss: 0.8264, Train Accuracy: 0.71% Validation Loss: 0.8685, Validation Accuracy: 0.70%

Epoch 4/20, Train Loss: 0.7337, Train Accuracy: 0.75% Validation Loss: 0.8363, Validation Accuracy: 0.71%

Epoch 5/20, Train Loss: 0.6686, Train Accuracy: 0.76% Validation Loss: 0.7935, Validation Accuracy: 0.73%  
Epoch 6/20, Train Loss: 0.6560, Train Accuracy: 0.77% Validation Loss: 0.7883, Validation Accuracy: 0.73%  
Epoch 7/20, Train Loss: 0.6099, Train Accuracy: 0.79% Validation Loss: 0.7984, Validation Accuracy: 0.73%  
Epoch 8/20, Train Loss: 0.5968, Train Accuracy: 0.79% Validation Loss: 0.8045, Validation Accuracy: 0.73%  
Epoch 9/20, Train Loss: 0.5437, Train Accuracy: 0.82% Validation Loss: 0.8088, Validation Accuracy: 0.73%  
Epoch 10/20, Train Loss: 0.5159, Train Accuracy: 0.83% Validation Loss: 0.8302, Validation Accuracy: 0.72%  
Epoch 11/20, Train Loss: 0.5146, Train Accuracy: 0.83% Validation Loss: 0.7458, Validation Accuracy: 0.75%  
Epoch 12/20, Train Loss: 0.5088, Train Accuracy: 0.82% Validation Loss: 0.8226, Validation Accuracy: 0.74%  
Epoch 13/20, Train Loss: 0.4967, Train Accuracy: 0.82% Validation Loss: 0.8047, Validation Accuracy: 0.74%  
Epoch 14/20, Train Loss: 0.4706, Train Accuracy: 0.84% Validation Loss: 0.7765, Validation Accuracy: 0.75%  
Epoch 15/20, Train Loss: 0.4416, Train Accuracy: 0.85% Validation Loss: 0.8384, Validation Accuracy: 0.73%  
Epoch 16/20, Train Loss: 0.4602, Train Accuracy: 0.85% Validation Loss: 0.7870, Validation Accuracy: 0.73%  
Epoch 17/20, Train Loss: 0.4167, Train Accuracy: 0.85% Validation Loss: 0.7979, Validation Accuracy: 0.74%  
Epoch 18/20, Train Loss: 0.4397, Train Accuracy: 0.85% Validation Loss: 0.8045, Validation Accuracy: 0.73%  
Epoch 19/20, Train Loss: 0.4051, Train Accuracy: 0.86% Validation Loss: 0.7393, Validation Accuracy: 0.74%  
Epoch 20/20, Train Loss: 0.3960, Train Accuracy: 0.87% Validation Loss: 0.7958, Validation Accuracy: 0.75%

### 2.2.3 Evaluating the model

```
1 from sklearn.metrics import confusion_matrix,
   ConfusionMatrixDisplay
2
3 all_labels = []
4 all_predictions = []
5
6 # Testing
7 model.eval()
8 test_loss = 0.0
9 correct_test = 0
10 total_test = 0
11
12 model.eval()
13 with torch.no_grad():
14     for inputs, labels in test_loader:
```

```

15         inputs, labels = inputs.to(device), labels.to(device)
16         outputs = model(inputs)
17         loss = criterion(outputs, labels)
18         test_loss += loss.item() * inputs.size(0)
19         _, predicted = outputs.max(1)
20         correct_test += (predicted == labels).sum().item()
21         total_test += labels.size(0)
22         all_labels.extend(labels.cpu().numpy())
23         all_predictions.extend(predicted.cpu().numpy())
24 test_accuracy = correct_test / total_test
25 test_loss /= total_test
26 print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {
    test_accuracy:.4f}")

```

Test Loss: 0.8071, Test Accuracy: 0.7274

```

1 import matplotlib.pyplot as plt
2
3 # Plot training and validation loss
4 plt.figure(figsize=(10, 5))
5 plt.plot(range(1, num_epochs + 1), train_losses, label="
    Training Loss")
6 plt.plot(range(1, num_epochs + 1), val_losses, label="
    Validation Loss")
7 plt.xlabel("Epochs")
8 plt.ylabel("Loss")
9 plt.title("Training and Validation Loss")
10 plt.legend()
11 plt.show()
12
13 # Plot training and validation accuracy
14 plt.figure(figsize=(10, 5))
15 plt.plot(range(1, num_epochs + 1), train_accuracies, label="
    Training Accuracy")
16 plt.plot(range(1, num_epochs + 1), val_accuracies, label="
    Validation Accuracy")
17 plt.xlabel("Epochs")
18 plt.ylabel("Accuracy")
19 plt.title("Training and Validation Accuracy")
20 plt.legend()
21 plt.show()

```




Figure 7: Training and Validation Loss




Figure 8: Training and Validation Accuracy

```

1 from sklearn.metrics import confusion_matrix,
  ConfusionMatrixDisplay
2 import matplotlib.pyplot as plt
3
4 # Generate confusion matrix
5 cm = confusion_matrix(all_labels, all_predictions)
6
7 # Create and plot the confusion matrix
8 disp = ConfusionMatrixDisplay(confusion_matrix=cm,
  display_labels=dataset.classes)
9 disp.plot(cmap='Blues', xticks_rotation=45)
10 plt.title('Confusion Matrix')
11 plt.show()

```




Figure 9: confusion matrix

## 2.3 DenseNet

### 2.3.1 Loading the pretrained model

```
1 model = models.densenet121(weights=models.DenseNet121_Weights.  
    IMAGENET1K_V1)  
2  
3 for param in model.parameters():  
4     param.requires_grad = False  
5  
6 num_features = model.classifier.in_features  
7 model.classifier = nn.Linear(num_features, num_classes)  
8  
9 nn.init.xavier_uniform_(model.classifier.weight)  
10 nn.init.zeros_(model.classifier.bias)  
11  
12 device = torch.device("cuda" if torch.cuda.is_available() else  
    "cpu")  
13 model = model.to(device)  
14  
15 criterion = nn.CrossEntropyLoss()  
16 optimizer = optim.Adam(model.classifier.parameters(), lr=  
    learning_rate)  
17  
18 print(f"Model is running on {device}")
```

Downloading: "https://download.pytorch.org/models/densenet121-a639ec97.pth"  
to /home/pasindupnk/.cache/torch/hub/checkpoints/densenet121-a639ec97.pth  
100%—— 30.8M/30.8M [00:03;00:00, 9.23MB/s] Model is running on cuda

### 2.3.2 Fine tuning the model

```
1 model = models.densenet121(weights=models.DenseNet121_Weights.  
    IMAGENET1K_V1)  
2  
3 for param in model.parameters():  
4     param.requires_grad = False  
5  
6 num_features = model.classifier.in_features  
7 model.classifier = nn.Linear(num_features, num_classes)  
8  
9 nn.init.xavier_uniform_(model.classifier.weight)  
10 nn.init.zeros_(model.classifier.bias)  
11  
12 device = torch.device("cuda" if torch.cuda.is_available() else  
    "cpu")  
13 model = model.to(device)  
14  
15 criterion = nn.CrossEntropyLoss()  
16 optimizer = optim.Adam(model.classifier.parameters(), lr=  
    learning_rate)  
17  
18 print(f"Model is running on {device}")
```

Epoch 1/20, Train Loss: 1.9297, Train Accuracy: 0.34, Validation Loss: 1.4091, Validation Accuracy: 0.49  
Epoch 2/20, Train Loss: 1.1414, Train Accuracy: 0.59, Validation Loss: 1.0602, Validation Accuracy: 0.62  
Epoch 3/20, Train Loss: 0.9154, Train Accuracy: 0.68, Validation Loss: 0.9550, Validation Accuracy: 0.67  
Epoch 4/20, Train Loss: 0.7823, Train Accuracy: 0.73, Validation Loss: 0.8637, Validation Accuracy: 0.69  
Epoch 5/20, Train Loss: 0.7062, Train Accuracy: 0.75, Validation Loss: 0.8227, Validation Accuracy: 0.71  
Epoch 6/20, Train Loss: 0.6511, Train Accuracy: 0.77, Validation Loss: 0.8039, Validation Accuracy: 0.71  
Epoch 7/20, Train Loss: 0.6012, Train Accuracy: 0.80, Validation Loss: 0.7795, Validation Accuracy: 0.72  
Epoch 8/20, Train Loss: 0.5862, Train Accuracy: 0.80, Validation Loss: 0.8148, Validation Accuracy: 0.72  
Epoch 9/20, Train Loss: 0.5587, Train Accuracy: 0.81, Validation Loss: 0.7501, Validation Accuracy: 0.75  
Epoch 10/20, Train Loss: 0.5173, Train Accuracy: 0.82, Validation Loss: 0.7492, Validation Accuracy: 0.73  
Epoch 11/20, Train Loss: 0.5029, Train Accuracy: 0.83, Validation Loss: 0.7497, Validation Accuracy: 0.73  
Epoch 12/20, Train Loss: 0.4957, Train Accuracy: 0.84, Validation Loss: 0.7355, Validation Accuracy: 0.73  
Epoch 13/20, Train Loss: 0.4605, Train Accuracy: 0.85, Validation Loss: 0.7160, Validation Accuracy: 0.75  
Epoch 14/20, Train Loss: 0.4649, Train Accuracy: 0.84, Validation Loss: 0.7437, Validation Accuracy: 0.74  
Epoch 15/20, Train Loss: 0.4401, Train Accuracy: 0.85, Validation Loss: 0.7675, Validation Accuracy: 0.73  
Epoch 16/20, Train Loss: 0.4303, Train Accuracy: 0.85, Validation Loss: 0.7416, Validation Accuracy: 0.74  
Epoch 17/20, Train Loss: 0.4217, Train Accuracy: 0.86, Validation Loss: 0.7365, Validation Accuracy: 0.74  
Epoch 18/20, Train Loss: 0.4163, Train Accuracy: 0.86, Validation Loss: 0.7271, Validation Accuracy: 0.74  
Epoch 19/20, Train Loss: 0.4251, Train Accuracy: 0.85, Validation Loss: 0.7258, Validation Accuracy: 0.75  
Epoch 20/20, Train Loss: 0.4052, Train Accuracy: 0.86, Validation Loss: 0.7356, Validation Accuracy: 0.74

### 2.3.3 Evaluating the model

```
1 # Testing
2 all_labels, all_predictions = [], []
3 model.eval()
4 test_loss, correct_test, total_test = 0.0, 0, 0
5
6 with torch.no_grad():
7     for inputs, labels in test_loader:
8         inputs, labels = inputs.to(device), labels.to(device)
9         outputs = model(inputs)
10        loss = criterion(outputs, labels)
11        test_loss += loss.item() * inputs.size(0)
12        _, predicted = outputs.max(1)
13        correct_test += (predicted == labels).sum().item()
14        total_test += labels.size(0)
15        all_labels.extend(labels.cpu().numpy())
16        all_predictions.extend(predicted.cpu().numpy())
17
18 test_accuracy = correct_test / total_test
19 test_loss /= total_test
20 print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {
    test_accuracy:.4f}")
```

**Test Loss: 0.7478, Test Accuracy: 0.7442**

```
1 model = models.resnet50(weights=models.ResNet50_Weights.
    IMAGENET1K_V1)
2
3 for param in model.parameters():
4     param.requires_grad = False
5
6 model.fc = nn.Linear(model.fc.in_features, num_classes)
7
8 nn.init.xavier_uniform_(model.fc.weight)
9 nn.init.zeros_(model.fc.bias)
10
11 device = torch.device("cuda" if torch.cuda.is_available() else
    "cpu")
12 model = model.to(device)
13
14 criterion = nn.CrossEntropyLoss()
15 optimizer = optim.Adam(model.fc.parameters(), lr=learning_rate
    )
16
17 print(f"Model is running on {device}")
```

**Model is running on cuda**



```

1 # Plot training and validation loss
2 plt.figure(figsize=(10, 5))
3 plt.plot(range(1, num_epochs + 1), train_losses, label="
4     Training Loss")
5 plt.plot(range(1, num_epochs + 1), val_losses, label="
6     Validation Loss")
7 plt.xlabel("Epochs")
8 plt.ylabel("Loss")
9 plt.title("Training and Validation Loss")
10 plt.legend()
11 plt.show()
12
13 # Plot training and validation accuracy
14 plt.figure(figsize=(10, 5))
15 plt.plot(range(1, num_epochs + 1), train_accuracies, label="
16     Training Accuracy")
17 plt.plot(range(1, num_epochs + 1), val_accuracies, label="
18     Validation Accuracy")
19 plt.xlabel("Epochs")
20 plt.ylabel("Accuracy")
21 plt.title("Training and Validation Accuracy")
22 plt.legend()
23 plt.show()

```




Figure 10: Training and Validation Loss




Figure 11: Training and Validation Accuracy

```
1 cm = confusion_matrix(all_labels, all_predictions)
2 disp = ConfusionMatrixDisplay(confusion_matrix=cm,
3                               display_labels=dataset.classes)
4 disp.plot(cmap='Blues', xticks_rotation=45)
5 plt.title('Confusion Matrix')
6 plt.show()
```




Figure 12: Confusion Matrix

## 2.4 Selecting pre-trained model or architecture

We have selected ResNet50 and DenseNet121, both well-known architectures pre-trained on ImageNet.

## 2.5 Fine-tuning the pre-trained model

The models are loaded using PyTorch's `torchvision.models` library, and their architectures are adapted using the `fine_tune_model` function to suit the Flowers-102 dataset. Specifically:

- **ResNet50:** The `fc` layer is modified to output the number of classes in the dataset.
- **DenseNet121:** The `classifier` layer is replaced to accommodate the new number of classes.

## 2.6 Train the fine-tuned model

The models are trained using the `train_model` function:

- **Optimizer:** Adam optimizer with a learning rate of 0.0001.
- **Learning Rate Scheduler:** Adjusts the learning rate every 7 epochs by a factor of 0.1.
- **Epochs:** 20.

Training splits (`train`, `val`, `test`) are organized into separate directories, and datasets are loaded using PyTorch's `ImageFolder` and `DataLoader` classes.

## 2.7 Training and validation loss values

Loss values are recorded for both training and validation splits and plotted using the `plot_loss` function.

## 2.8 Evaluate the fine-tuned model

The `evaluate_model` function computes the test accuracy using the testing split. It utilizes a simple loop to accumulate correct predictions and calculate the percentage accuracy.

## 2.9 Compare the test accuracy

To evaluate and compare the effectiveness of pre-trained models like ResNet or DenseNet with a custom CNN, several metrics should be considered, including accuracy, loss, confusion matrices, inference time, and generalization capabilities. Pre-trained models often excel due to their sophisticated architectures and transfer learning advantages, leveraging features learned from large datasets like ImageNet. If they outperform the custom CNN, it highlights their ability to adapt well to new tasks with limited data. Conversely, if the custom CNN performs comparably or better, it suggests that the model is well-designed and fine-tuned for the specific dataset and task. Visualization through accuracy and loss curves, along with confusion matrices, helps analyze classification performance and identify misclassification patterns. Ultimately, discussing factors like dataset size, complexity, and model parameters provides insights into why one approach outperforms the other, offering a deeper understanding of model behaviour.

## 2.10 Discussion trade-offs, advantages, and limitations

### 2.10.1 Custom Model

#### Advantages

- **Lightweight:** Custom models are smaller in size, which makes them faster to train and easier to deploy on resource-constrained devices.

- **Full Customizability:** They can be tailored to the specific requirements of the dataset and task, allowing for optimized performance when designed appropriately.
- **Simpler Architecture:** Easier to understand, debug, and modify compared to complex pre-trained models.

### Limitations

- **Limited Generalization:** Custom models may underperform on complex datasets, particularly when data is scarce, as they lack prior knowledge from extensive pretraining.
- **Requires More Training Data:** Without access to pre-trained weights, custom models need significantly larger datasets to learn high-level features effectively.
- **Shallow Learning:** May struggle with extracting intricate patterns due to limited depth and absence of pretraining.

### 2.10.2 Pretrained Model

#### Advantages

- **Faster Convergence:** pretrained models start with learned weights, enabling them to recognize general features like edges and shapes, reducing training time.
- **Better Generalization:** Leveraging knowledge from large datasets (e.g., ImageNet) improves performance on smaller datasets, particularly when data is limited.
- **Robustness:** pretrained weights act as a form of regularization, mitigating overfitting in cases of limited data.
- **Complex Pattern Recognition:** Handles intricate patterns and domain-specific features more effectively than custom models.

#### Limitations

- **Computational Cost:** Fine-tuning pretrained models is computationally intensive and requires more resources for both training and inference.
- **Domain Gap Challenges:** When the target dataset differs significantly from the source dataset (e.g., ImageNet vs. medical imaging), adaptation becomes challenging.
- **Hyperparameter Sensitivity:** Fine-tuning requires careful adjustment of hyperparameters, such as learning rate, to avoid overfitting or underfitting.
- **Larger Model Size:** pre-trained models consume more memory and storage, making them less suitable for lightweight applications.

## Trade-offs and Comparison

- **Performance vs. Complexity:** pre-trained models generally outperform custom models due to their depth and pretraining, but they come with increased computational costs and complexity.
- **Flexibility vs. Generalization:** Custom models are easier to customize and may be ideal for specific, simple tasks, while pre-trained models are better suited for complex or high-stakes tasks requiring robust generalization.
- **Training Data Requirements:** Custom models need substantial labelled data to perform well, whereas pre-trained models thrive even with limited data through transfer learning.

Pretrained models are typically the better choice for tasks with limited data or complex patterns, as they offer faster convergence, higher accuracy, and better generalization. However, custom models are advantageous when computational resources are constrained, the task is straightforward, or there's a need for a lightweight and tailored solution. The choice depends on the dataset size, task complexity, and resource availability.

## Appendix

You can find the project repository on GitHub:

[Image Classification using CNN](#)